



NATIONAL UNIVERSITY
of Computer & Emerging Sciences

DATA MINING PROJECT

Group Members

- Syed Muhammad Faizan Abbas Shah Kazmi (211-5675)
- Muhammad Talal Saleem (211-5682)
- Muhammad Huzair Amjad (211-6213)

Instructor's Name

Dr. Esha Tu Razia Babar

Project Report: Asteroid Hazard Classification

Abstract

This report presents the methodology, findings, and insights gained from a data mining project focused on classifying asteroids into hazardous and non-hazardous categories and predicting their hazard potential based on various features. Leveraging a dataset sourced from NeoWs (Near Earth Object Web Service), the project aims to develop a predictive model capable of determining whether an asteroid poses a hazard to Earth, while identifying the key features contributing to its hazardous status. The report details the steps undertaken in data preprocessing, feature engineering, model selection, evaluation metrics, hazard prediction, and visualization. The expected outcome is a robust classification model that accurately predicts asteroid hazard potential, thereby enhancing our ability to assess the risks associated with near-Earth asteroids.

Introduction

Asteroids pose a potential threat to Earth, with the potential for catastrophic consequences if not properly monitored and mitigated. Understanding the hazard potential of near-Earth asteroids is crucial for planetary defense and safeguarding human civilization. This project addresses the need for accurate classification and prediction of asteroid hazard potential by employing data mining techniques on a comprehensive dataset sourced from NeoWs. The development of a robust predictive model holds promise for enhancing our ability to assess and manage the risks posed by near-Earth asteroids.

Problem and Dataset Description

This project aims to leverage machine learning techniques to build a robust predictive model that effectively categorizes asteroids into hazardous and non-hazardous classes. By harnessing a dataset obtained from NeoWs, which encapsulates pertinent details about asteroids like Absolute Magnitude, Estimated Diameter, Close Approach Date, Relative Velocity, and a binary indicator denoting Hazardous status, we endeavor to explore and

analyze the intrinsic relationships between these features and the likelihood of an asteroid being hazardous. Through comprehensive data preprocessing, feature engineering, and model selection, our goal is to develop a predictive framework capable of accurately discerning potentially dangerous asteroids from benign ones. Ultimately, this endeavor not only enhances our understanding of asteroid behavior and characteristics but also contributes to the ongoing efforts in planetary defense and space exploration.

Dataset Description

The dataset encompasses a wide range of features relevant to asteroid classification and hazard prediction. Each entry includes details such as Neo Reference ID, Absolute Magnitude, Estimated Diameter, Orbital Parameters, and a binary indicator for Hazardous status. A comprehensive list of features and their descriptions is provided in the project proposal.

Methodology

The methodology involves several key steps:

Data Preprocessing:

Handling missing values, normalizing numerical features, and encoding categorical variables to prepare the dataset for modeling.

- The code starts by importing necessary libraries (**pandas**, **numpy**, **matplotlib.pyplot**, **seaborn**) and loading the dataset using **pd.read_csv()**.

	Neo Reference ID	Name	Absolute Magnitude	Est Dia in KM(min)	Est Dia in KM(max)	Est Dia in M(min)	Est Dia in M(max)	Est Dia in Miles(min)	Est Dia in Miles(max)	Est Dia in Feet(min)	...	Asc Node Longitude	Orbital Period	Perihelion Distance	Perihelion Arg	Aphelic Dis
0	3703080	3703080	21.6	0.127220	0.284472	127.219879	284.472297	0.079051	0.176763	417.388066	...	314.373913	609.599786	0.808259	57.257470	2.00576
1	3723955	3723955	21.3	0.146068	0.326618	146.067964	326.617897	0.090762	0.202951	479.225620	...	136.717242	425.869294	0.718200	313.091975	1.49735
2	2446882	2446882	20.3	0.231502	0.517654	231.502122	517.654482	0.143849	0.321655	759.521423	...	259.475979	643.580228	0.950791	248.415038	1.96885
3	3092506	3092506	27.4	0.008801	0.019681	8.801465	19.680675	0.005469	0.012229	28.876199	...	57.173266	514.082140	0.983902	18.707701	1.52790
4	3514799	3514799	21.6	0.127220	0.284472	127.219879	284.472297	0.079051	0.176763	417.388066	...	84.629307	495.597821	0.967687	158.263596	1.48354

- It checks for missing values and drops columns with significant missing values (`data.isnull().sum()` and `data.drop()`).

Missing Values:	
Neo Reference ID	0
Name	0
Absolute Magnitude	0
Est Dia in KM(min)	0
Est Dia in KM(max)	0
Est Dia in M(min)	0
Est Dia in M(max)	0
Est Dia in Miles(min)	0
Est Dia in Miles(max)	0
Est Dia in Feet(min)	0
Est Dia in Feet(max)	0
Close Approach Date	0
Epoch Date Close Approach	0
Relative Velocity km per sec	0
Relative Velocity km per hr	0
Miles per hour	0
Miss Dist.(Astronomical)	0
Miss Dist.(lunar)	0
Miss Dist.(kilometers)	0
Miss Dist.(miles)	0
Orbiting Body	0
Orbit ID	0
Orbit Determination Date	0
Orbit Uncertainty	0
Minimum Orbit Intersection	0
Jupiter Tisserand Invariant	0

- It converts categorical variables to numerical using label encoding (`LabelEncoder`) for non-numeric columns.

```
Equinox
J2000      4687
Name: count, dtype: int64
```

```
Orbiting Body
Earth      4687
Name: count, dtype: int64
```

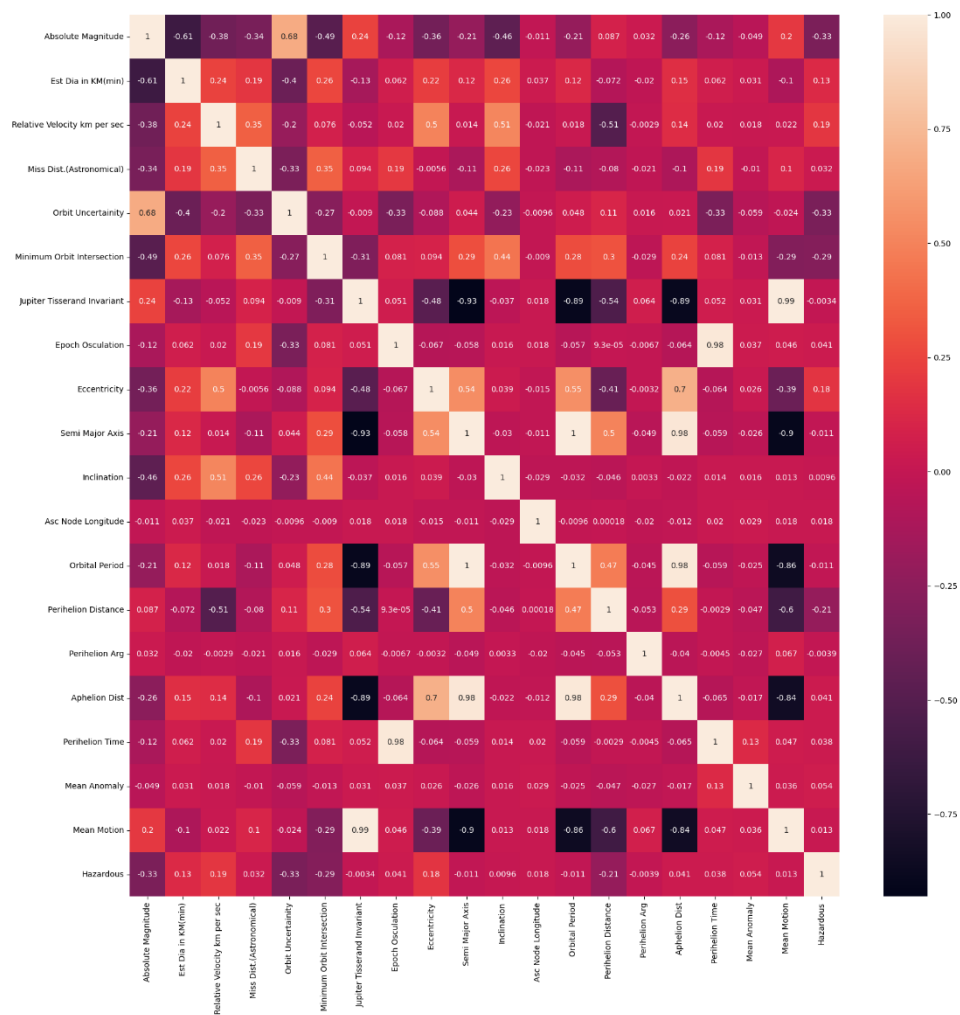
- It drops additional columns that are not required for analysis.
Before:

```
4      3514799      3514799
5 rows x 40 columns
```

After:

```
4      21.6      0.127220      0.28
5 rows x 34 columns
```

- Visualizes the correlation matrix using a heatmap (`sns.heatmap()`).



Feature Engineering:

Extracting relevant features and performing dimensionality reduction if necessary to enhance model performance.

- Variance thresholding, SelectKBest (**Chi-squared** and **ANOVA F-value**), Recursive Feature Elimination (**RFE**), and Principal Component Analysis (**PCA**) are used for feature selection.

```
[ ]
Principal Components:
Absolute Magnitude  Est Dia in KM(min)  Est Dia in KM(max)  \
0 -2.460032e-12      1.755439e-13      3.925281e-13
1  3.618838e-08      -2.563721e-09      -5.732656e-09
2 -2.697132e-05      2.383581e-06      5.329849e-06
3 -4.948506e-04      1.189563e-04      2.659944e-04
4  9.712952e-05      2.372703e-06      5.305526e-06

Est Dia in M(min)  Est Dia in M(max)  Est Dia in Miles(min)  \
0  1.755439e-10      3.925281e-10      1.090779e-13
1 -2.563722e-06      -5.732656e-06      -1.593022e-09
2  2.383581e-03      5.329849e-03      1.481088e-06
3  1.189563e-01      2.659944e-01      7.391599e-05
4  2.372703e-03      5.305526e-03      1.474329e-06

Est Dia in Miles(max)  Est Dia in Feet(min)  Est Dia in Feet(max)  \
0  2.439056e-13      5.759315e-10      1.287822e-09
1 -3.562106e-09      -8.411160e-06      -1.880793e-05
2  3.311813e-06      7.820147e-03      1.748638e-02
3  1.652812e-04      3.902766e-01      8.726849e-01
4  3.296700e-06      7.784460e-03      1.740658e-02

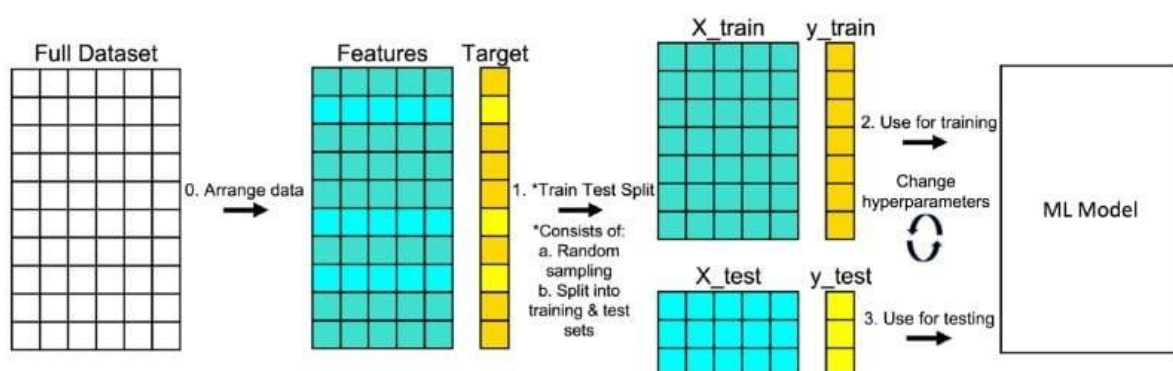
Epoch Date Close Approach ... Inclination Asc Node Longitude \
0 -1.000000e+00 ... 3.716359e-12 1.008034e-11
1 -1.876563e-05 ... -1.070453e-07 1.061231e-07
2  8.750895e-09 ... 1.714409e-04 -5.503398e-05
3  8.000013e-10 ... 4.741170e-04 -1.523673e-02
```

- These techniques help in selecting the most relevant features for the models.

Model Training & Selection:

Exploring various classification techniques including Decision Trees, Support Vector Machines, Neural Networks, and ensemble methods to identify the most suitable model for the task.

- It uses libraries (**sklearn,xgboost,lightgbm,catboost,tensorflow,keras**) for machine learning.
- Split** your data into training and testing sets, then train your chosen model on the training data



- The code trains several machine learning **models** such as Decision Trees, Random Forest, K-Nearest Neighbors (KNN), Naive Bayes, Support Vector Machines (SVM), XGBoost, Neural Networks, and Gradient Boosting.

- It uses techniques like **SMOTE** for handling class imbalance.

Before:

```
Class Distribution:
Hazardous
0      3932
1       755
Name: count, dtype: int64
```

After:

```
Class Distribution after SMOTE:
Hazardous
1      3932
0      3932
Name: count, dtype: int64
```

- **Model evaluation** metrics like accuracy, precision, recall, F1-score, specificity, and confusion matrix are calculated for each model.
- Hyperparameter tuning is performed using **GridSearchCV** to find the best parameters for models like Random Forest, SVM, XGBoost, Gradient Boosting, etc.

Evaluation Metrics:

Assessing model performance using standard classification metrics such as:

- **Accuracy** (Calculate the accuracy of your model on the test set to measure its overall performance)
- **Precision** (Calculate the ratio of true positive (**TP**) predictions to the total number of positive predictions made by the classifier)
- **Recall** (Calculate the ratio of true positive (**TP**) predictions to the total number of actual positive instances in the dataset)
- **F1-score** (It provides a balance between precision and recall)
- **Confusion Matrix** (Analyze the confusion matrix to understand the model's predictions in terms of true positives (**TP**), true negatives (**TN**), false positives (**FP**), and false negatives (**FN**))

- **ROC curve and AUC** (Plot the Receiver Operating Characteristic (**ROC**) curve and calculate the Area Under the Curve (**AUC**) score to assess the model's ability to discriminate between hazardous and non-hazardous asteroids)

By considering the class imbalance in the dataset.

Evaluation Metrics for Decision Tree Classifier:

```

Accuracy: 0.9974570883661793
Confusion Matrix:
[[797  4]
 [ 0 772]]
Precision: 0.9948453608247423
Recall: 1.0
Sensitivity: 1.0
Specificity: 0.9950062421972534
F1 Score: 0.9974160206718347
Classification Report:

```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	801
1	0.99	1.00	1.00	772
accuracy			1.00	1573
macro avg	1.00	1.00	1.00	1573
weighted avg	1.00	1.00	1.00	1573

Evaluation Metrics for Random Forest Classifier:

```

Accuracy: 0.9987285441830897
Confusion Matrix:
[[799  2]
 [ 0 772]]
Precision: 0.9974160206718347
Recall: 1.0
Sensitivity: 1.0
Specificity: 0.9975031210986267
F1 Score: 0.9987063389391979
Classification Report:

```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	801
1	1.00	1.00	1.00	772
accuracy			1.00	1573
macro avg	1.00	1.00	1.00	1573
weighted avg	1.00	1.00	1.00	1573

Evaluation Metrics for KNN:

```
Accuracy: 0.7889383343928799
Confusion Matrix:
[[538 263]
 [ 69 703]]
Precision: 0.727743271221532
Recall: 0.9106217616580311
Sensitivity: 0.9106217616580311
Specificity: 0.6716604244694132
F1 Score: 0.8089758342922899
Classification Report:
      precision    recall  f1-score   support

     0       0.89       0.67       0.76       801
     1       0.73       0.91       0.81       772

 accuracy          0.79       1573
 macro avg       0.81       0.79       0.79       1573
 weighted avg    0.81       0.79       0.79       1573
```

Evaluation Metrics for Naïve Bayes:

```
Accuracy: 0.891290527654164
Confusion Matrix:
[[686 115]
 [ 56 716]]
Precision: 0.8616125150421179
Recall: 0.927461139896373
Sensitivity: 0.927461139896373
Specificity: 0.8564294631710362
F1 Score: 0.8933250155957578
Classification Report:
      precision    recall  f1-score   support

     0       0.92       0.86       0.89       801
     1       0.86       0.93       0.89       772

 accuracy          0.89       1573
 macro avg       0.89       0.89       0.89       1573
 weighted avg    0.89       0.89       0.89       1573
```

Evaluation Metrics for SVM:

```
Accuracy: 0.9828353464717101
Confusion Matrix:
[[782  19]
 [  8 764]]
Classification Report:
              precision    recall  f1-score   support

     0       0.99       0.98       0.98        801
     1       0.98       0.99       0.98        772

 accuracy          0.98
 macro avg         0.98       0.98       0.98
weighted avg         0.98       0.98       0.98

Precision: 0.9757343550446999
Recall: 0.9896373056994818
Sensitivity: 0.9896373056994818
F1 Score: 0.982636655948553
Specificity: 0.9762796504369539
```

Evaluation Metrics for XGBOOST:

```
Accuracy: 0.9980928162746344
Confusion Matrix:
[[798   3]
 [  0 772]]
Precision: 0.9961290322580645
Recall: 1.0
Sensitivity: 1.0
Specificity: 0.9962546816479401
F1 Score: 0.9980607627666451
Classification Report:
              precision    recall  f1-score   support

     0       1.00       1.00       1.00        801
     1       1.00       1.00       1.00        772

 accuracy          1.00
 macro avg         1.00       1.00       1.00
weighted avg         1.00       1.00       1.00
```

Evaluation Metrics for Neural Network:

```
Accuracy: 0.9847425301970757
Confusion Matrix:
[[792  9]
 [ 15 757]]
Precision: 0.9882506527415144
Recall: 0.9805699481865285
Sensitivity: 0.9805699481865285
F1 Score: 0.9843953185955786
Specificity: 0.9887640449438202
Classification Report:
              precision    recall  f1-score   support

     0           0.98         0.99         0.99         801
     1           0.99         0.98         0.98         772

   accuracy          0.98
  macro avg          0.98
 weighted avg          0.98
```

Evaluation Metrics for Gradient Boost:

```
Accuracy: 0.9980928162746344
Confusion Matrix:
[[798  3]
 [  0 772]]
Precision: 0.9961290322580645
Recall: 1.0
Sensitivity: 1.0
Specificity: 0.9962546816479401
F1 Score: 0.9980607627666451
Classification Report:
              precision    recall  f1-score   support

     0           1.00         1.00         1.00         801
     1           1.00         1.00         1.00         772

   accuracy          0.99
  macro avg          1.00
 weighted avg          1.00
```

When using GridCV:

Evaluation Metrics for Random Forest:

```
Accuracy: 0.9972979286813797 Parameters: {'max_depth': 20, 'min_samples_leaf': 4, 'min_samples_split': 10, 'n_estimators': 100}
[ ] Accuracy: 0.9976158937052271 Parameters: {'max_depth': 20, 'min_samples_leaf': 4, 'min_samples_split': 10, 'n_estimators': 200}
Accuracy: 0.9972979286813797 Parameters: {'max_depth': 20, 'min_samples_leaf': 4, 'min_samples_split': 10, 'n_estimators': 300}

Best Parameters: {'max_depth': None, 'min_samples_leaf': 1, 'min_samples_split': 10, 'n_estimators': 200}

Best Accuracy: 0.9980928162746344
Precision: 0.9974126778783958
Recall: 0.9987046632124352
Sensitivity: 0.9987046632124352
Specificity: 0.9975031210986267
F1 Score: 0.9980582524271844

Confusion Matrix:
[[799  2]
 [ 1 771]]

Classification Report:
      precision    recall  f1-score   support

     0       1.00      1.00      1.00        801
     1       1.00      1.00      1.00        772

   accuracy          1.00      1.00      1.00      1573
  macro avg          1.00      1.00      1.00      1573
 weighted avg          1.00      1.00      1.00      1573
```

Evaluation Metrics for Decision Tree Classifier:

```
Accuracy: 0.9974570883661793
Precision: 0.9961240310077519
Recall: 0.9987046632124352
Sensitivity: 0.9987046632124352
Specificity: 0.9962546816479401
F1 Score: 0.9974126778783958

Confusion Matrix:
[[798  3]
 [ 1 771]]

Classification Report:
      precision    recall  f1-score   support

     0       1.00      1.00      1.00        801
     1       1.00      1.00      1.00        772

   accuracy          1.00      1.00      1.00      1573
  macro avg          1.00      1.00      1.00      1573
 weighted avg          1.00      1.00      1.00      1573
```

Evaluation Metrics for KNN:

```
[ ] Best Parameters: { 'algorithm': 'auto', 'k_neighbors': 3, 'p': 1, 'weights': 'distance' }

Accuracy: 0.87730451366815
Precision: 0.8150163220892275
Recall: 0.9702072538860104
Sensitivity: 0.9702072538860104
Specificity: 0.787765293383271
F1 Score: 0.8858663512714371
Confusion matrix:
[[631 170]
 [ 23 749]]
Classification Report:

```

	precision	recall	f1-score	support
0	0.96	0.79	0.87	801
1	0.82	0.97	0.89	772
accuracy			0.88	1573
macro avg	0.89	0.88	0.88	1573
weighted avg	0.89	0.88	0.88	1573

Evaluation Metrics for Naïve Bayes:

```
Accuracy: 0.891290527654164
Precision: 0.8616125150421179
Recall: 0.927461139896373
Sensitivity: 0.927461139896373
Specificity: 0.8564294631710362
F1 Score: 0.8933250155957578

Confusion Matrix:
[[686 115]
 [ 56 716]]
Classification Report:

```

	precision	recall	f1-score	support
0	0.92	0.86	0.89	801
1	0.86	0.93	0.89	772
accuracy			0.89	1573
macro avg	0.89	0.89	0.89	1573
weighted avg	0.89	0.89	0.89	1573

Evaluation Metrics for SVM:

```
Best Parameters: {'C': 20, 'gamma': 'auto', 'kernel': 'rbf'}

Accuracy: 0.9866497139224412
Precision: 0.982028241335045
Recall: 0.9909326424870466
Sensitivity: 0.9909326424870466
Specificity: 0.982521847690387
F1 Score: 0.9864603481624759

Confusion Matrix:
[[787  14]
 [  7 765]]

Classification Report:
      precision    recall  f1-score   support

     0       0.99       0.98       0.99         801
     1       0.98       0.99       0.99         772

   accuracy          0.99
  macro avg          0.99
 weighted avg          0.99
```

Evaluation Metrics for XGBOOST:

```
Best Accuracy: 0.9974570883661793
Precision: 0.9961240310077519
Recall: 0.9987046632124352
Sensitivity: 0.9987046632124352
Specificity: 0.9962546816479401
F1 Score: 0.9974126778783958

Confusion Matrix:
[[798   3]
 [  1 771]]

Classification Report:
      precision    recall  f1-score   support

     0       1.00       1.00       1.00         801
     1       1.00       1.00       1.00         772

   accuracy          1.00
  macro avg          1.00
 weighted avg          1.00
```

Evaluation Metrics for Neural Networks:

```
Best Parameters: {'activation': 'relu', 'alpha': 0.01, 'hidden_layer_sizes': 100,

Accuracy: 0.9942784488239034
Precision: 0.990990990990991
Recall: 0.9974093264248705
Sensitivity: 0.9974093264248705
Specificity: 0.9912609238451935
F1 Score: 0.9941897998708844

Confusion Matrix:
[[794  7]
 [ 2 770]]

Classification Report:
      precision    recall  f1-score   support

      0         1.00      0.99      0.99         801
      1         0.99      1.00      0.99         772

   accuracy          0.99          0.99          0.99         1573
  macro avg          0.99          0.99          0.99         1573
 weighted avg          0.99          0.99          0.99         1573
```

Evaluation Metrics for Gradient Boost:

```
Best Parameters: {'learning_rate': 0.05, 'max_depth': 6, 'max_features': 'log2',

Accuracy: 0.9993642720915448
Precision: 0.9987063389391979
Recall: 1.0
Sensitivity: 1.0
Specificity: 0.9987515605493134
F1 Score: 0.9993527508090615

Confusion Matrix:
[[800  1]
 [ 0 772]]

Classification Report:
      precision    recall  f1-score   support

      0         1.00      1.00      1.00         801
      1         1.00      1.00      1.00         772

   accuracy          1.00          1.00          1.00         1573
  macro avg          1.00          1.00          1.00         1573
 weighted avg          1.00          1.00          1.00         1573
```

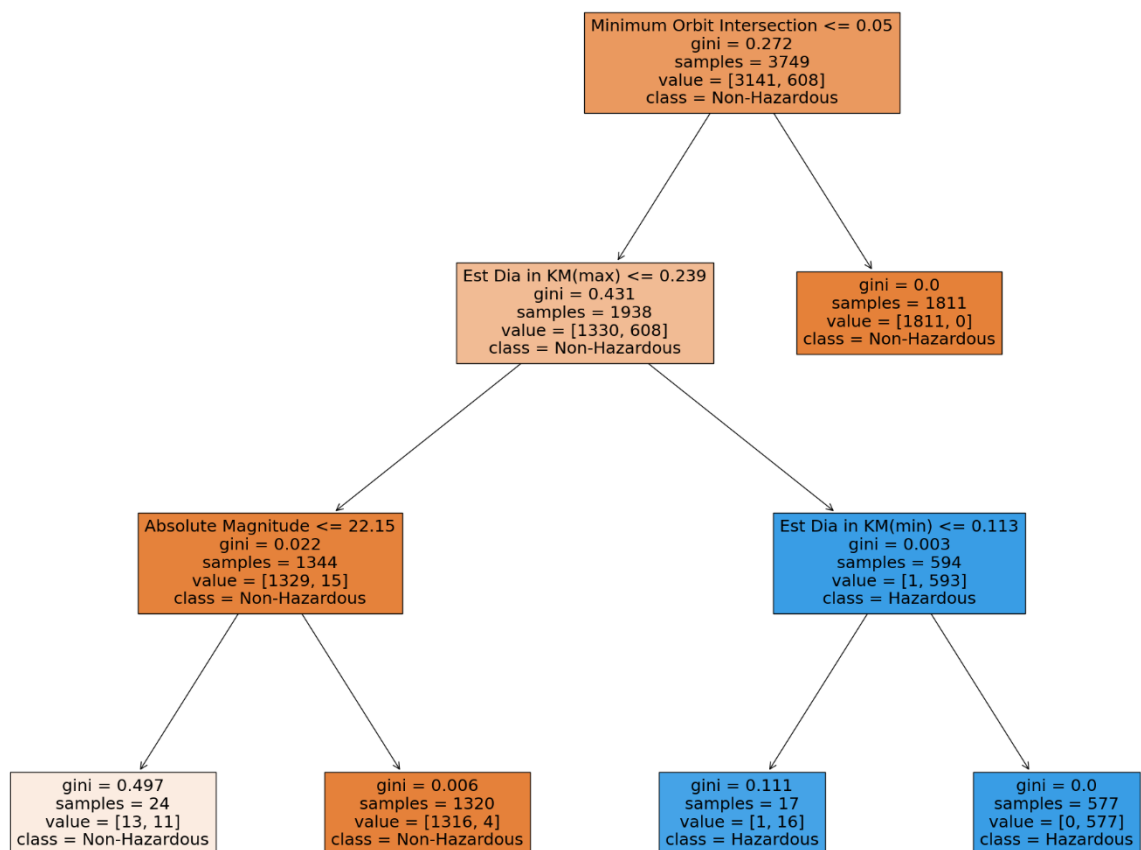
Hazard Prediction:

Utilizing feature importance techniques to identify key factors contributing to asteroid hazard potential.

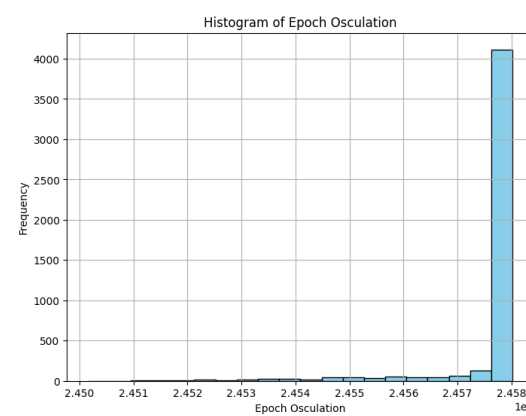
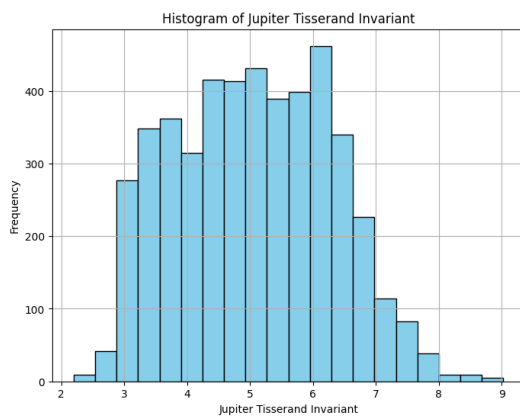
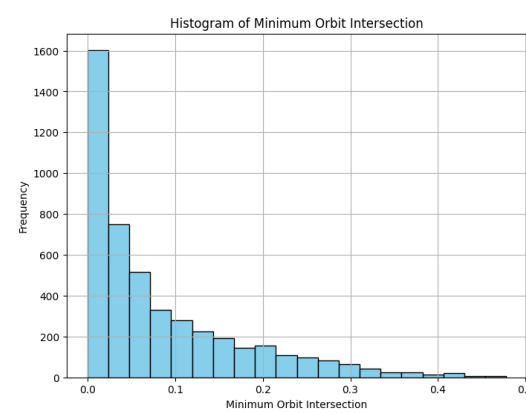
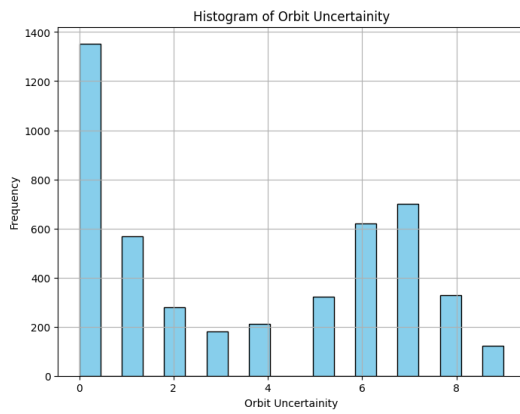
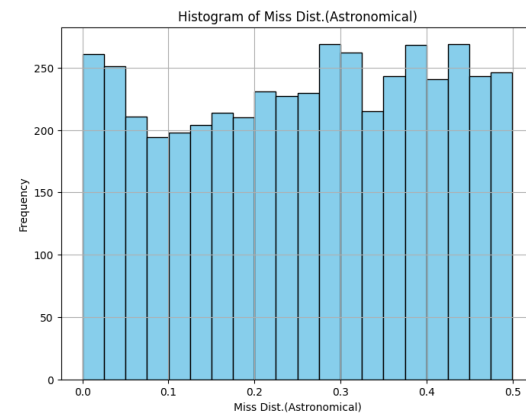
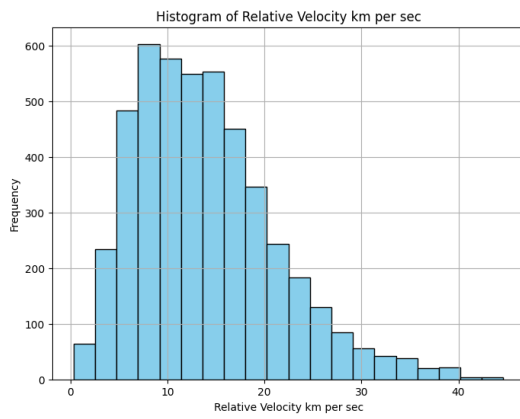
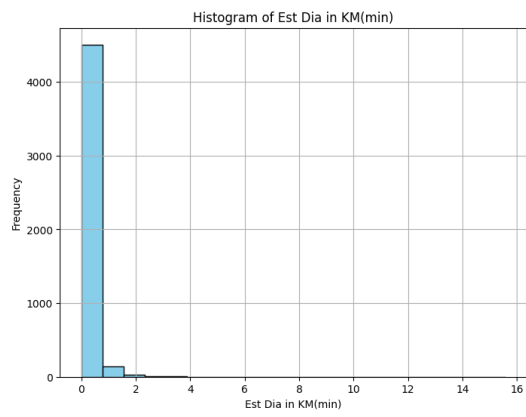
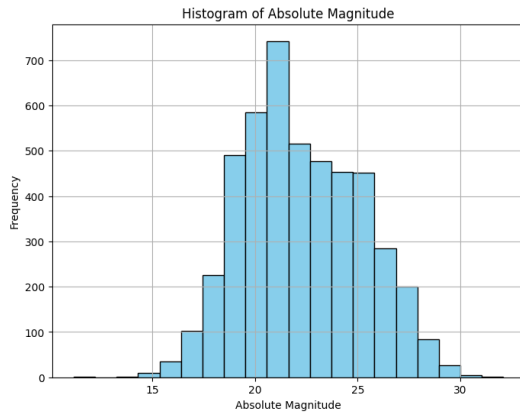
Visualization:

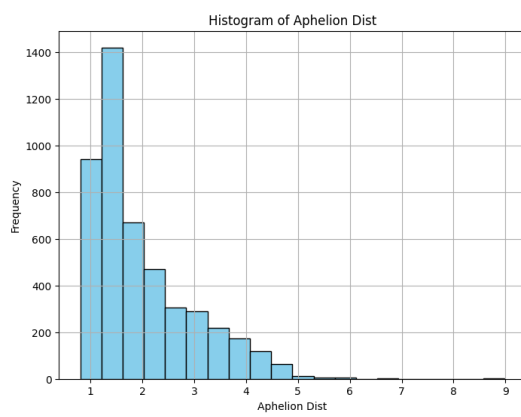
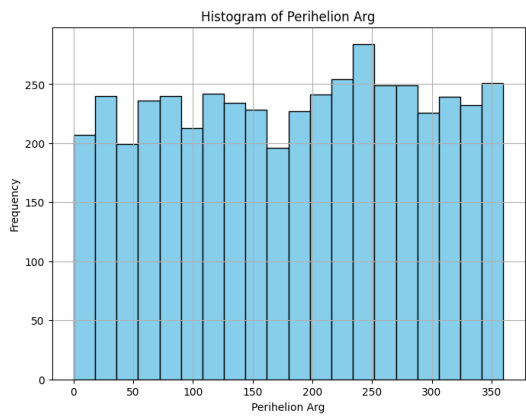
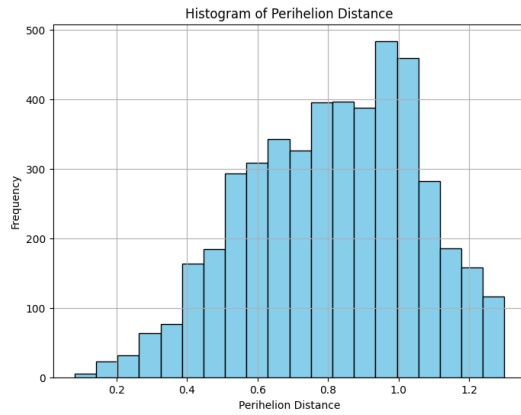
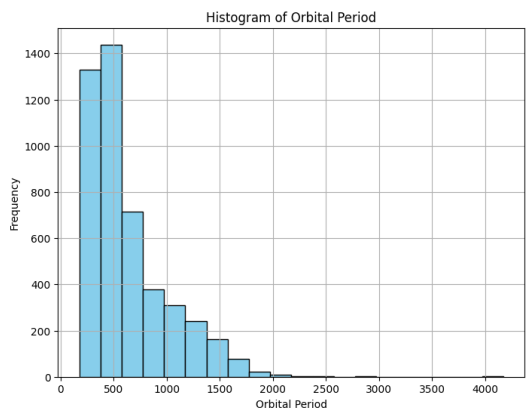
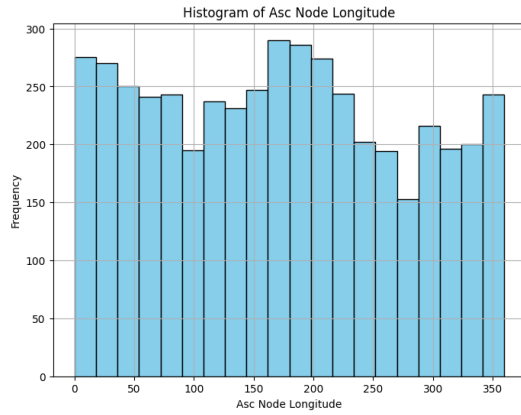
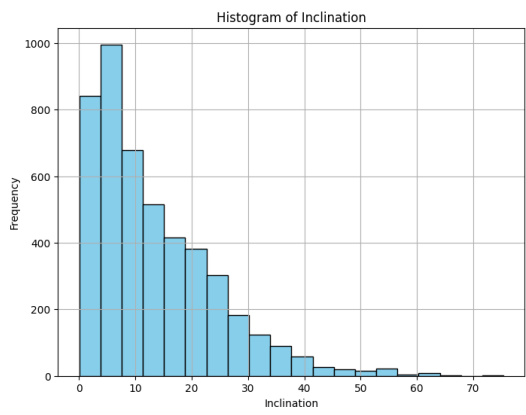
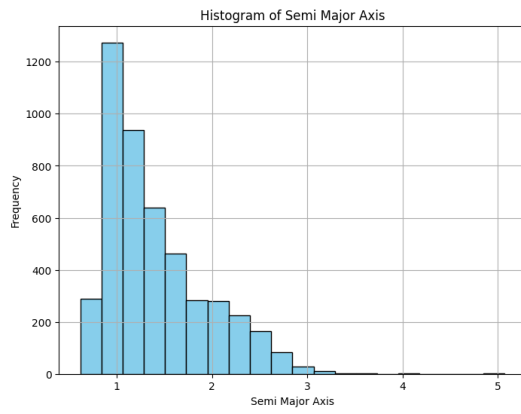
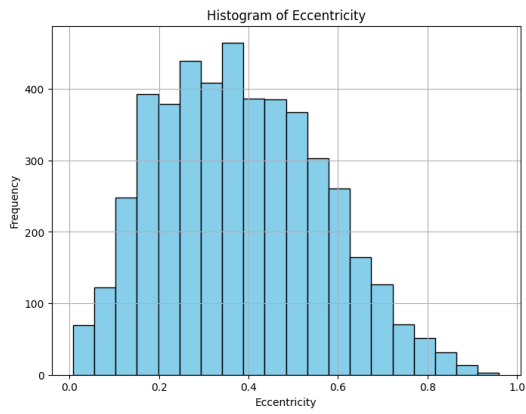
Visualizing the dataset and model predictions using plots and graphs to gain insights into the data distribution and model performance.

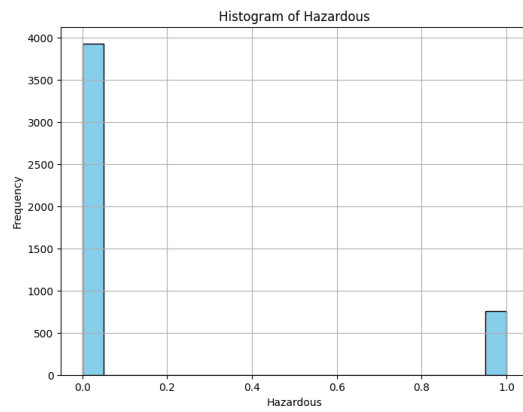
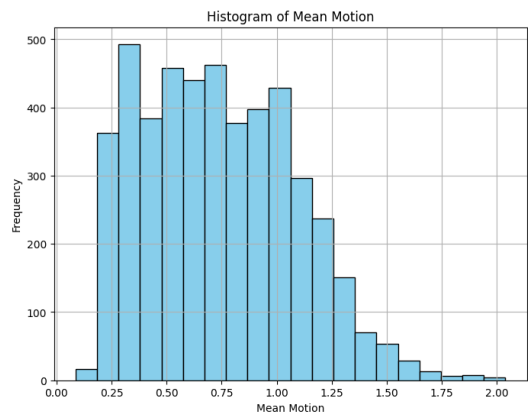
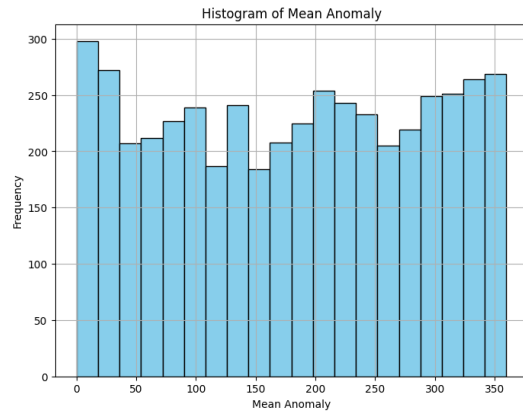
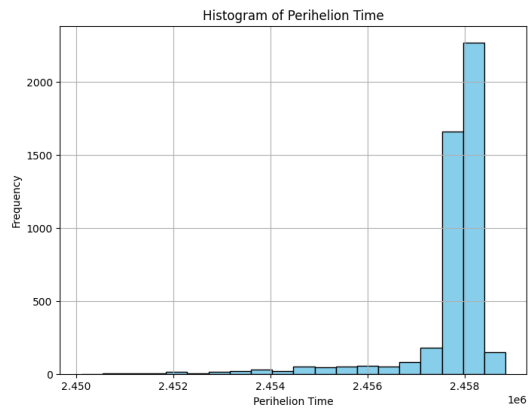
- **Decision Tree** visualization is performed to understand the decision-making process of the model.



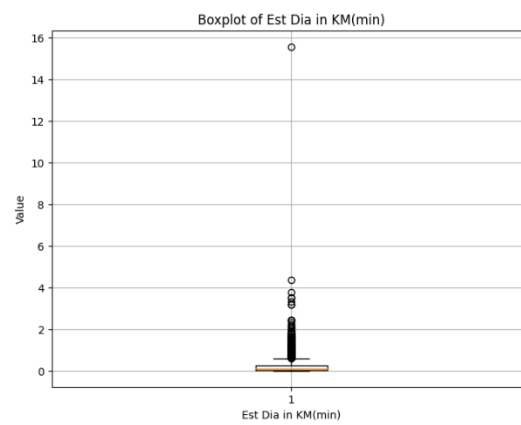
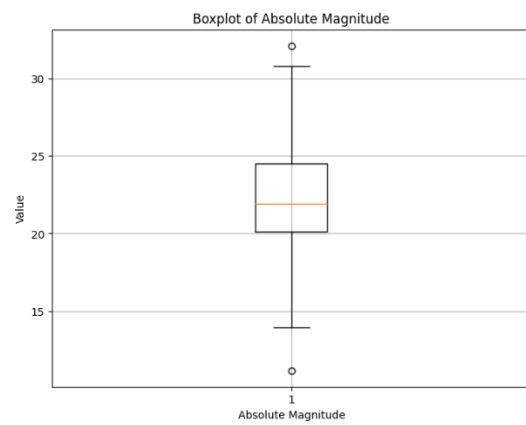
■ Histograms:

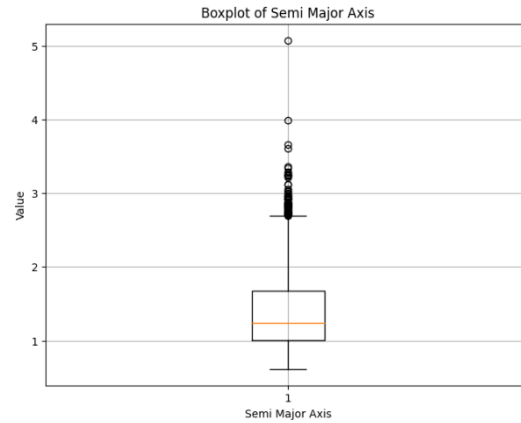
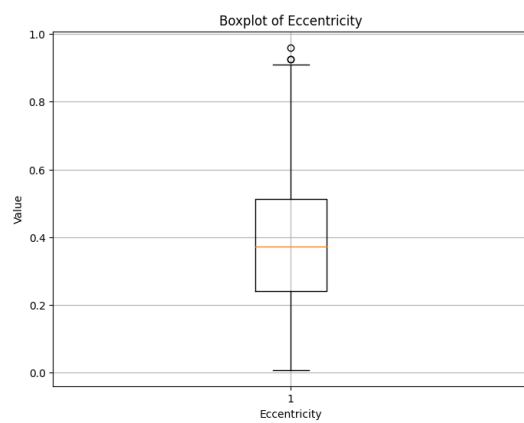
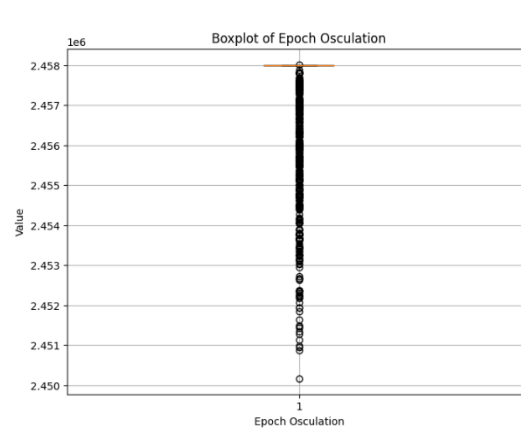
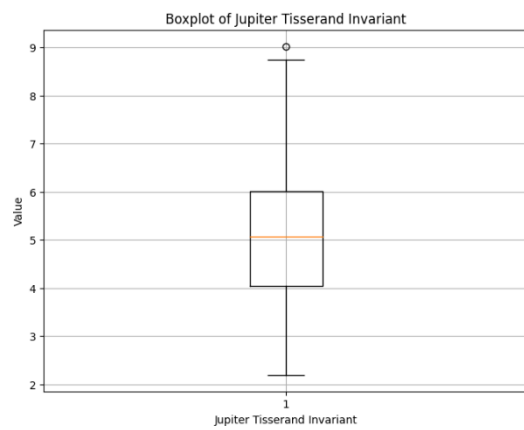
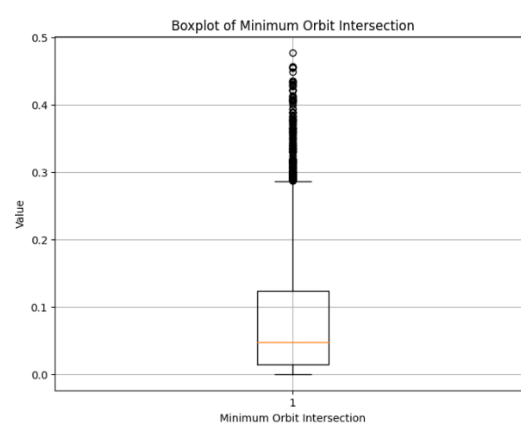
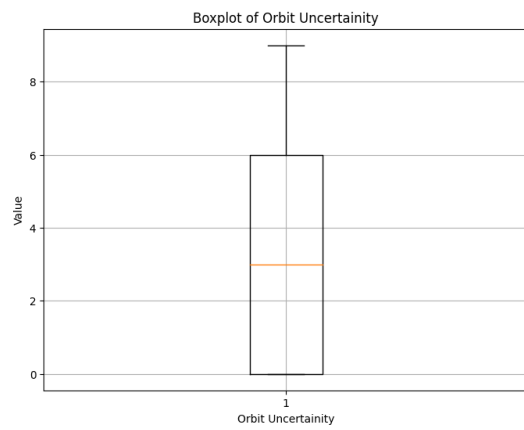
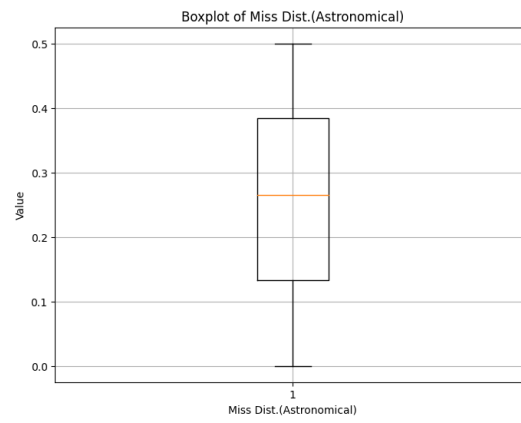
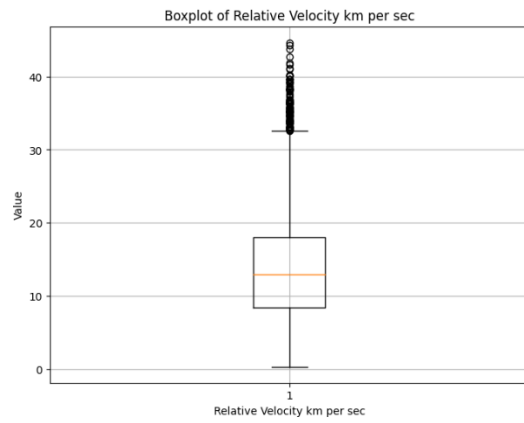


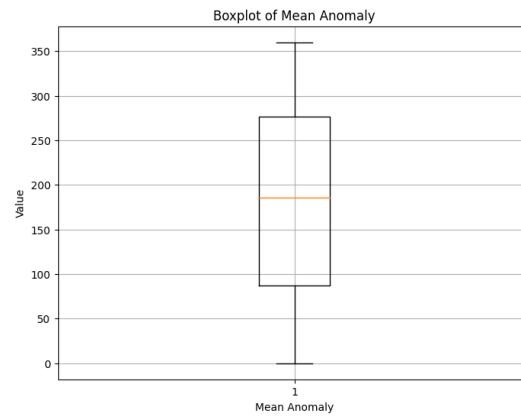
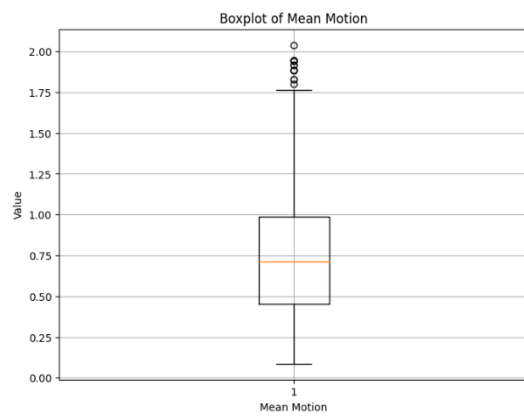
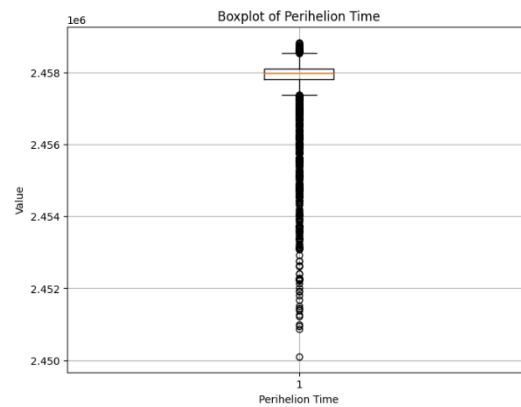
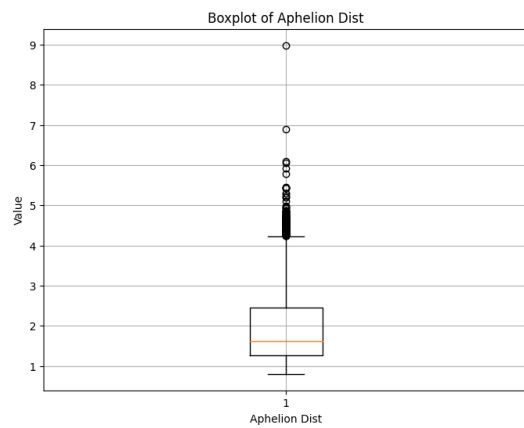
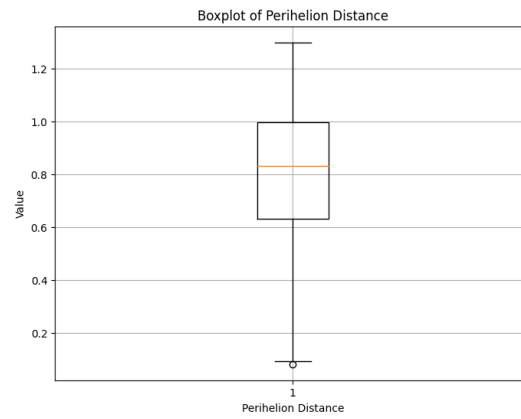
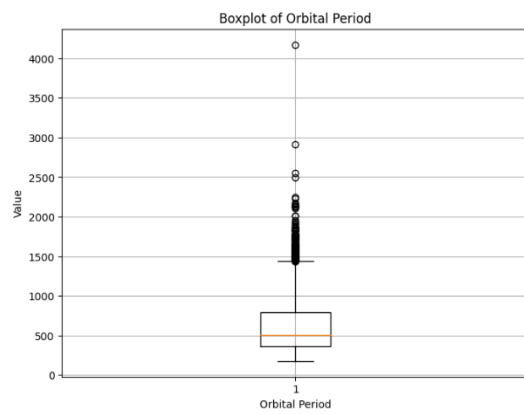
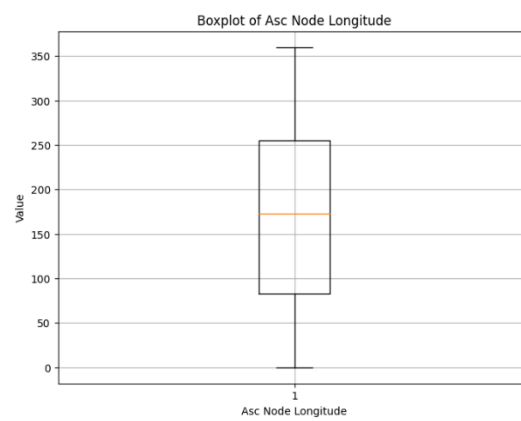
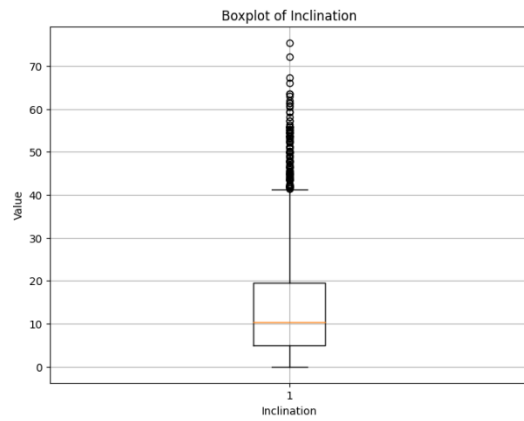




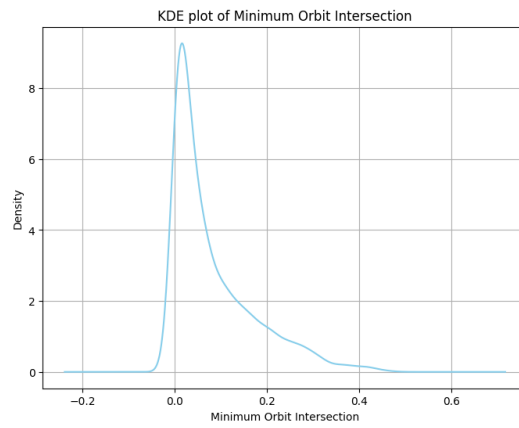
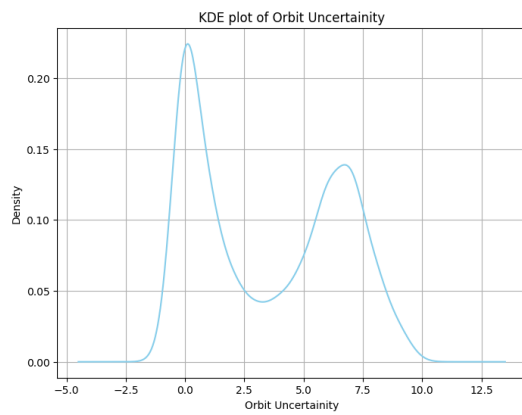
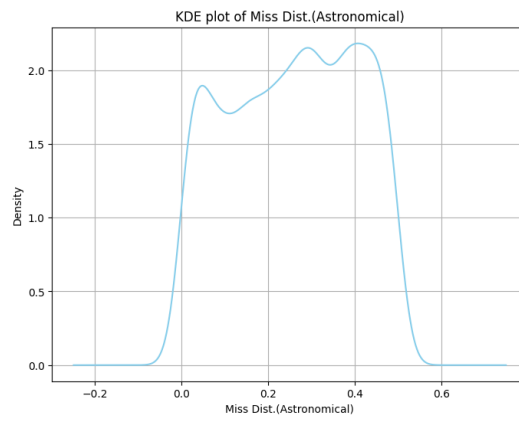
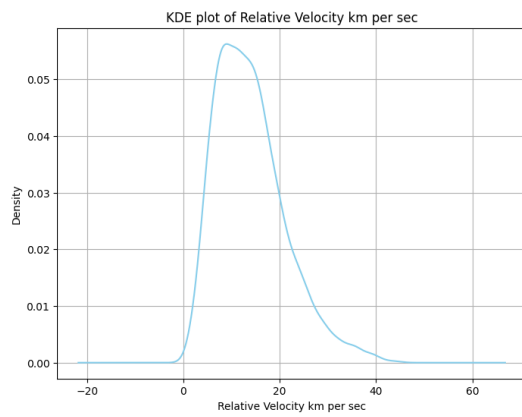
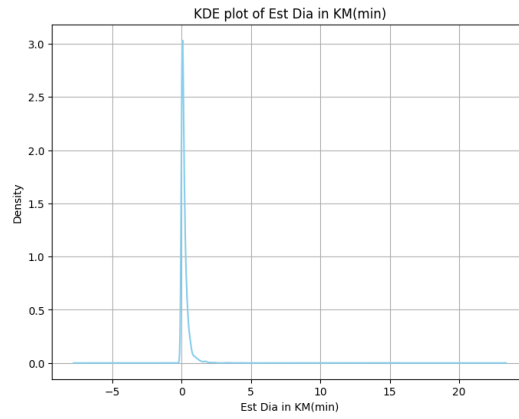
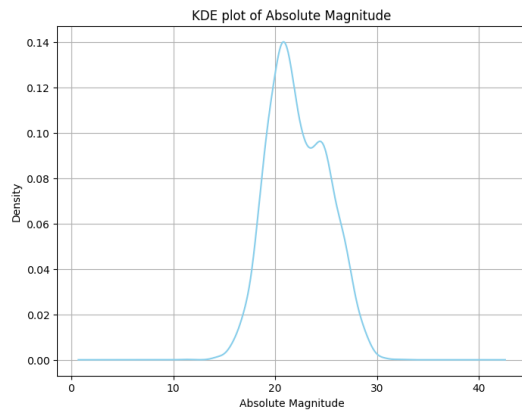
Box Plots:

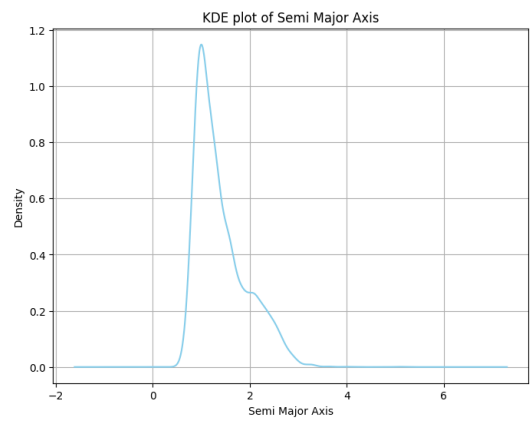
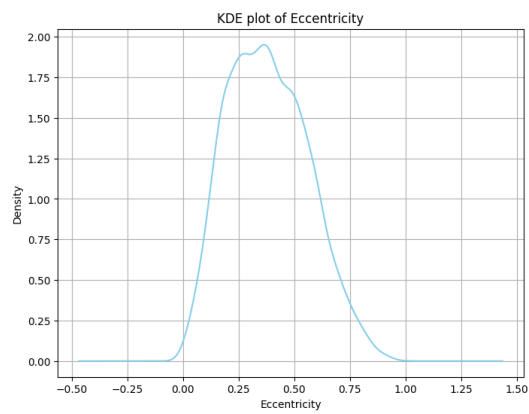
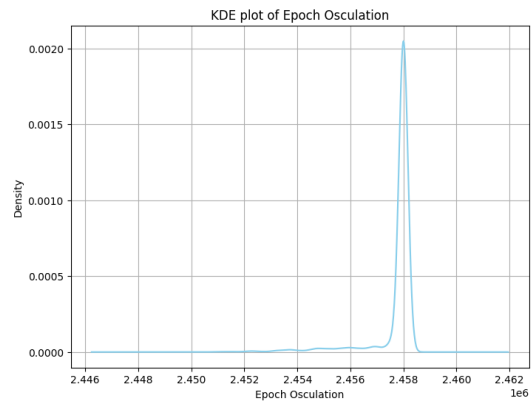
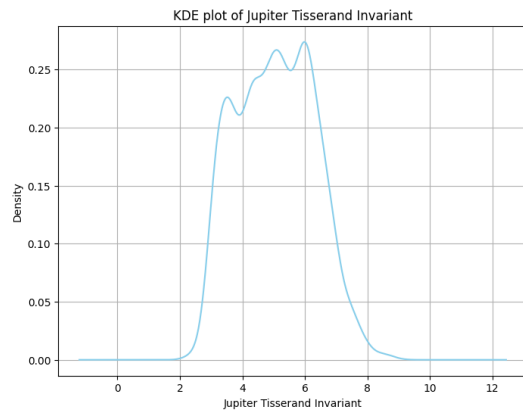




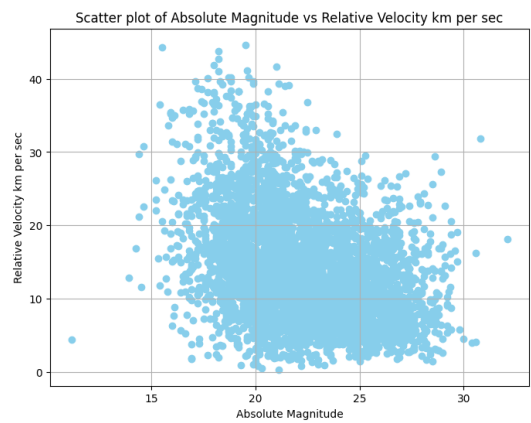
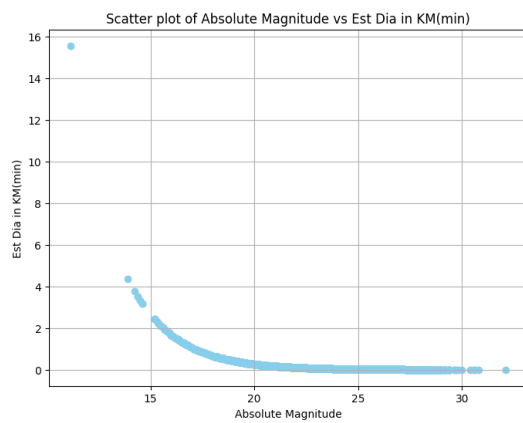


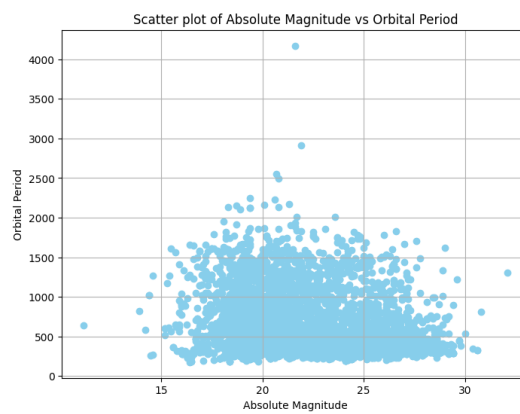
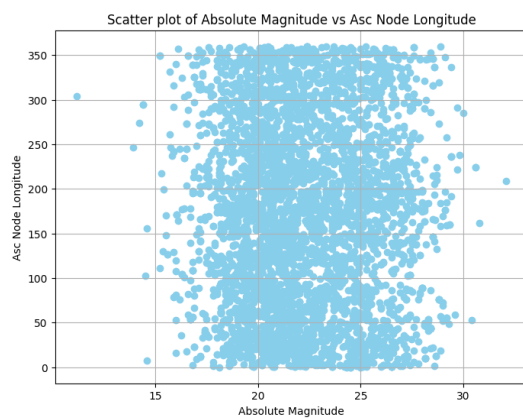
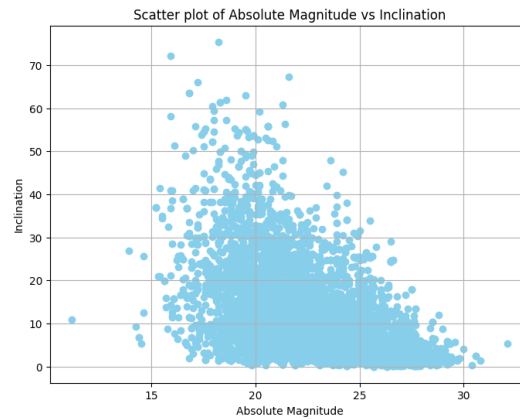
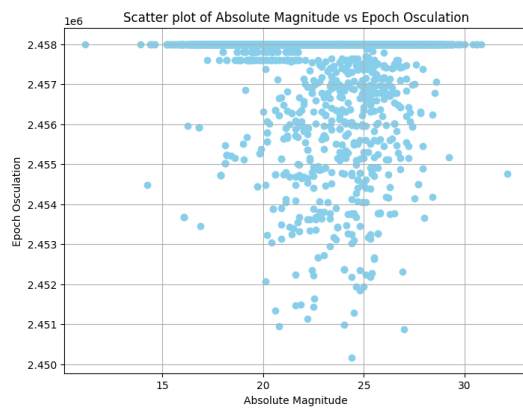
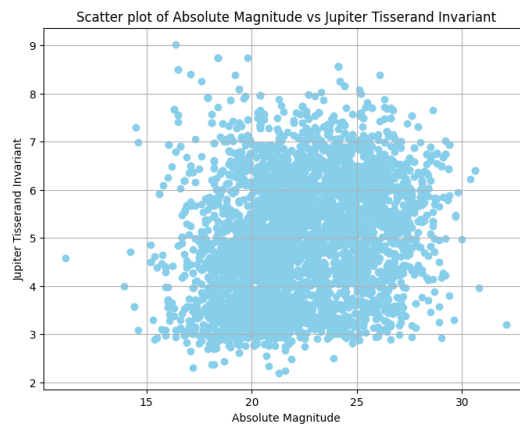
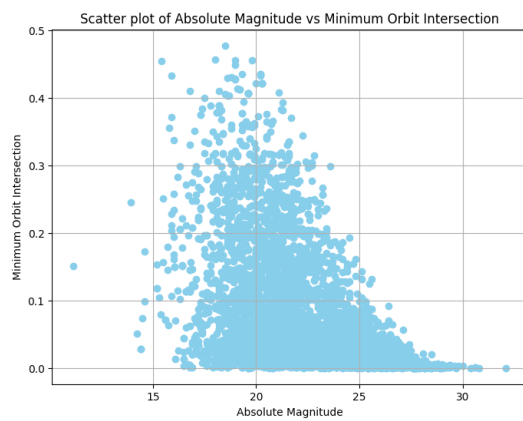
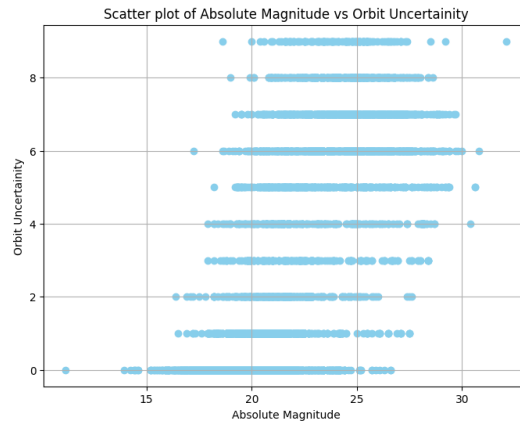
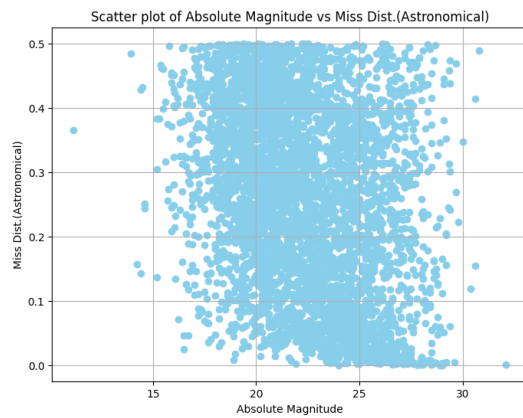
■ Kernel Density Estimate (KDE):

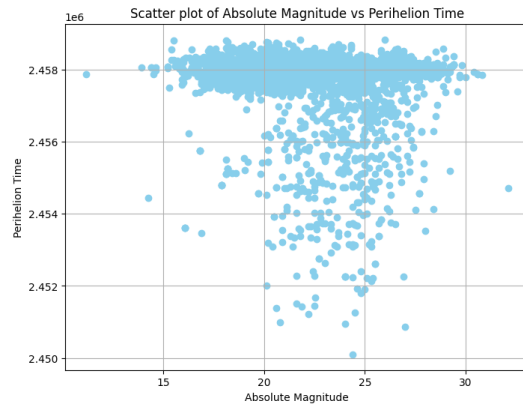
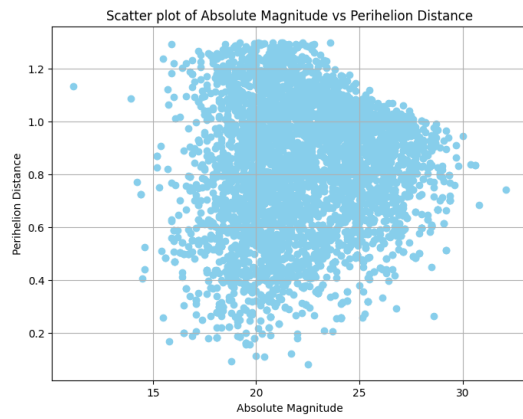




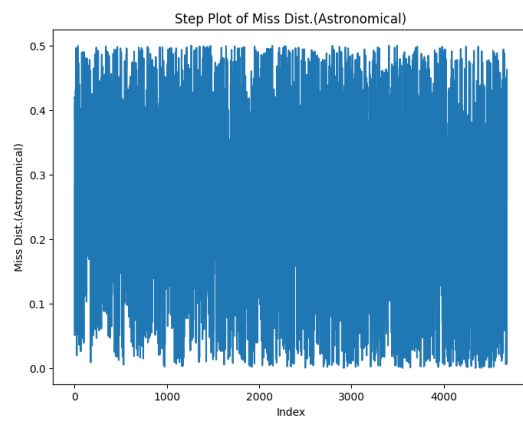
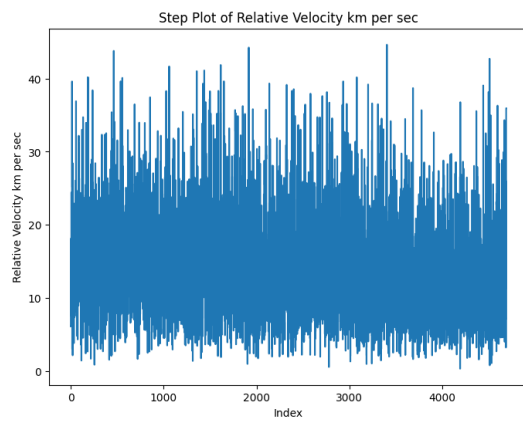
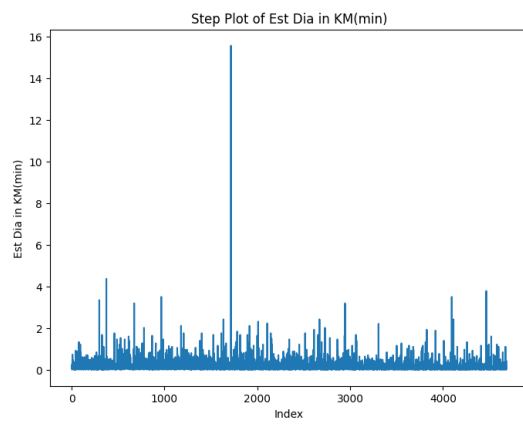
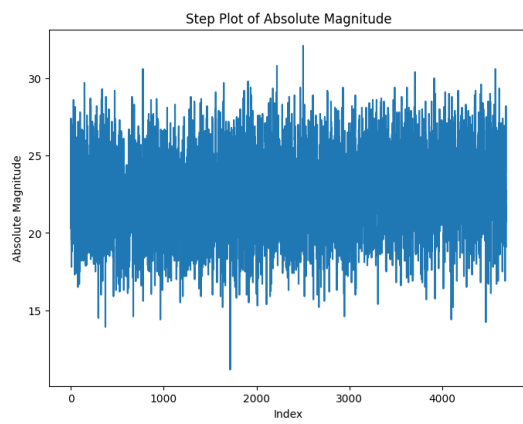
■ Scatter Plot:

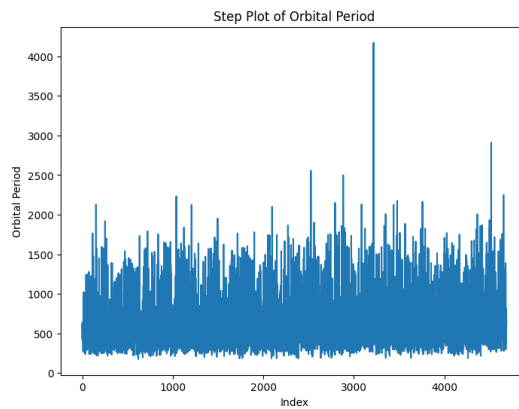
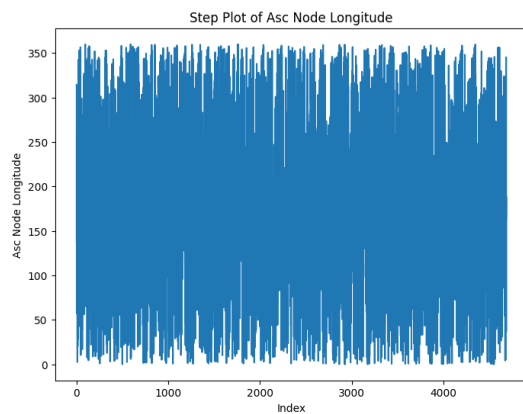
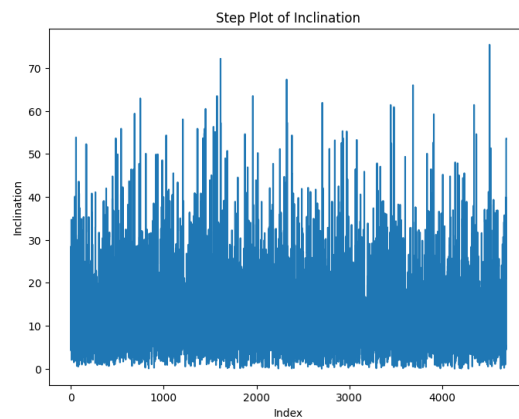
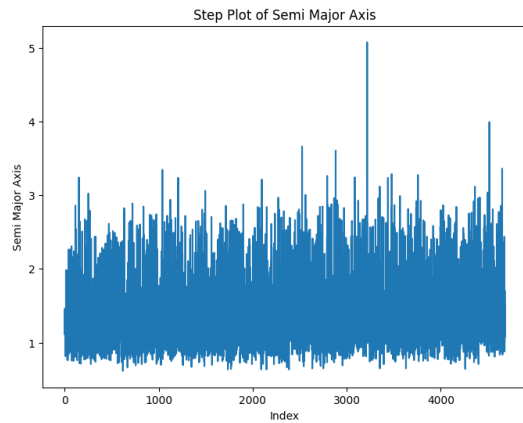
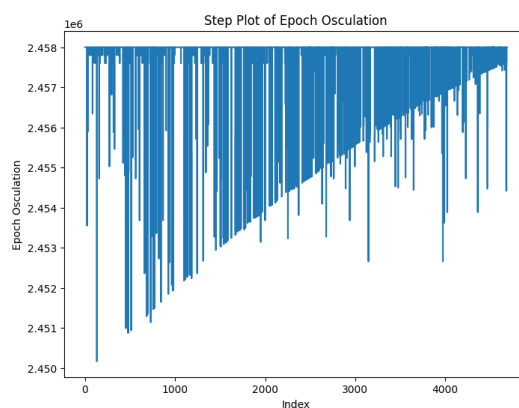
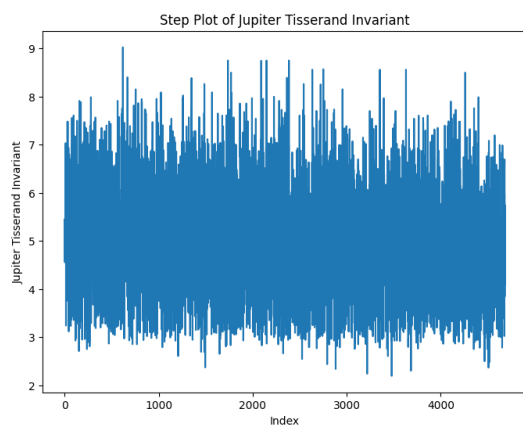
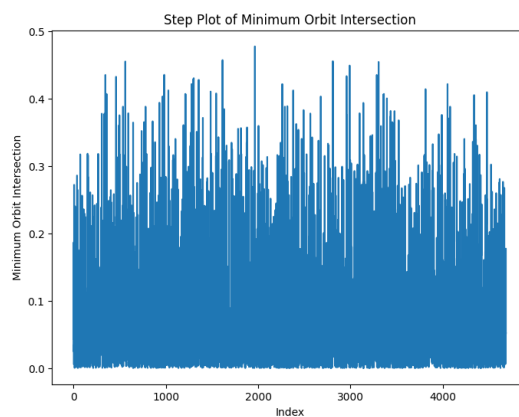
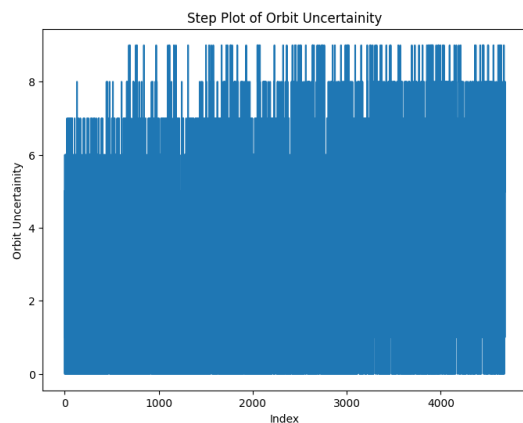


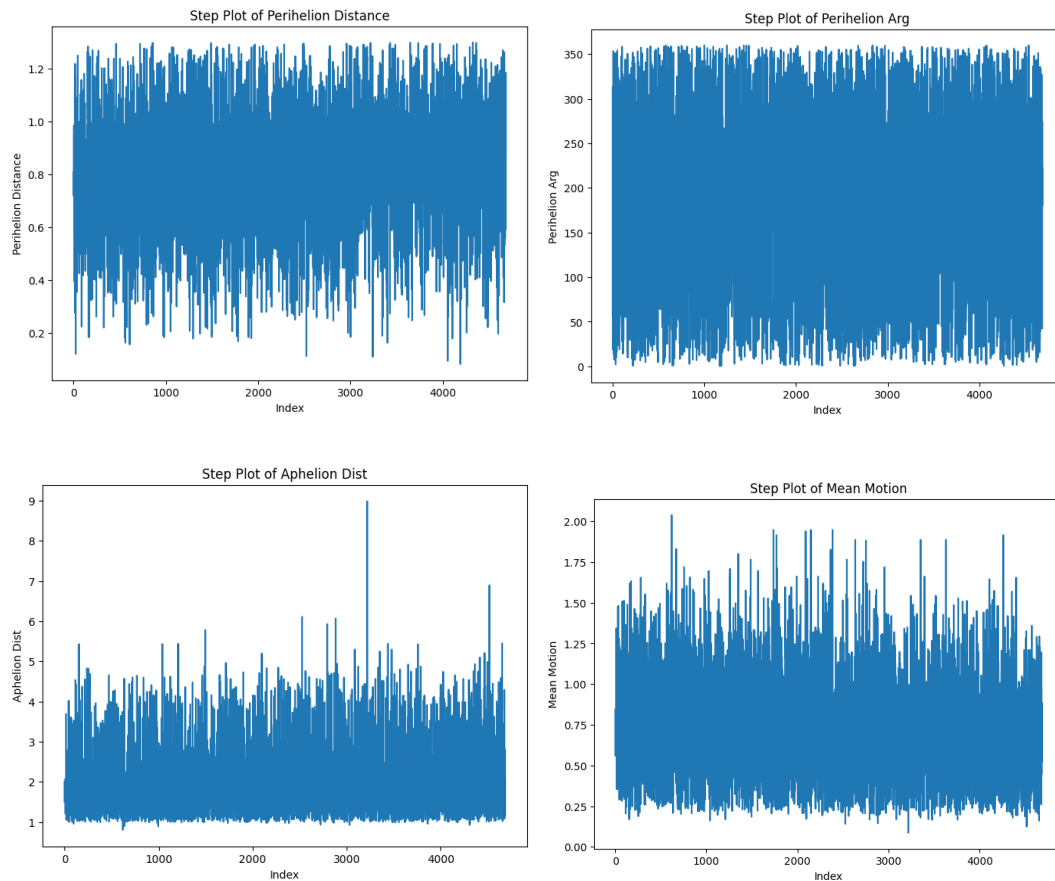




Bar Graphs:







Software Tools:

Jupyter Notebook or Google Colab will be utilized for data exploration, preprocessing, model development, and evaluation. Libraries such as **Pandas**, **NumPy**, **Scikit-learn**, and **Matplotlib** will facilitate data manipulation, modeling, and visualization.

DM_Proj.ipynb

File Edit View Insert Runtime Tools Help Last edited on 10 May

+ Code + Text

```
[ ] import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

[ ] data=pd.read_csv("nasa.csv")

[ ] data.head()
```

	Neo Reference ID	Name	Absolute Magnitude	Est Dia in KM(min)	Est Dia in KM(max)	Est Dia in M(min)	Est Dia in M(max)	Est Dia in Miles(min)	Est Dia in Miles(max)	Est Dia in Feet(min)	...	Asc Node Longitude	Orbital Period	Perihelion Distance	Perihelion Arg	Aphelion Dis
0	3703080	3703080	21.6	0.127220	0.284472	127.219879	284.472297	0.079051	0.176763	417.388066	...	314.373913	609.599786	0.808259	57.257470	2.00576
1	3723955	3723955	21.3	0.146068	0.326618	146.067964	326.617897	0.090762	0.202951	479.225620	...	136.717242	425.869294	0.718200	313.091975	1.49735
2	2446862	2446862	20.3	0.231502	0.517654	231.502122	517.654482	0.143849	0.321655	759.521423	...	259.475979	643.580228	0.950791	248.415038	1.96688
3	3092506	3092506	27.4	0.008801	0.019681	8.801465	19.880675	0.005469	0.012229	28.876199	...	57.173266	514.082140	0.983902	18.707701	1.52796
4	3514799	3514799	21.6	0.127220	0.284472	127.219879	284.472297	0.079051	0.176763	417.388066	...	84.629307	495.597821	0.967687	158.263596	1.48354

5 rows x 40 columns

Expected Results and Evaluation Techniques:

The expected outcome is a classification model that accurately predicts whether an asteroid is hazardous or not, based on its features. Model evaluation will be conducted using standard classification metrics such as accuracy, precision, recall, and F1-score. Additionally, ROC-AUC curve analysis will assess the model's ability to discriminate between hazardous and non-hazardous asteroids.

Preliminary Results and Dataset Exploration:

The dataset will be explored through Exploratory Data Analysis (EDA) to understand feature distribution, detect missing values or outliers, and identify potential challenges. Visualization techniques will be employed to gain insights into feature-target relationships. Preprocessing steps such as handling missing values, encoding categorical variables, and feature scaling will be performed as needed before model training.

Outline of the Work-to-Do:

Data preprocessing: Handling missing values, normalizing features, and encoding categorical variables.

Feature engineering: Extracting relevant features and performing dimensionality reduction if necessary.

Model training: Experimenting with different classification algorithms and hyperparameters.

Model evaluation: Assessing model performance using appropriate evaluation metrics.

Feature importance analysis: Identifying key features influencing asteroid hazard potential.

Visualization: Visualizing data distributions, model predictions, and feature importance to facilitate interpretation.

Conclusion

This project stands as a pivotal response to the pressing necessity for precise classification and predictive modeling concerning the potential hazards posed by **asteroids**, employing sophisticated data mining methodologies. Through the development of a resilient **classification model** and the discernment of pivotal features, the overarching goal is to significantly elevate our capacity to evaluate and effectively mitigate the risks entwined with near-Earth asteroids. By delineating and comprehensively analyzing various asteroid properties and orbital characteristics, this endeavor strives to furnish stakeholders with invaluable insights into the nature and severity of **potential threats**, thereby empowering informed decision-making and proactive risk management strategies. Moreover, the iterative nature of this project ensures that further refinement and validation of the model will be rigorously pursued, leveraging the insights gleaned throughout its course to continually enhance the **accuracy, reliability, and applicability** of the developed predictive framework. In doing so, this initiative not only contributes to the advancement of scientific understanding in the field of asteroid hazard assessment but also holds the potential to safeguard lives, infrastructure, and planetary well-being through proactive and informed risk mitigation measures.

References

- [1] NeoWs (Near Earth Object Web Service). Follow: <https://catalog.data.gov/dataset/>
- [2] Pandas documentation. Follow: <https://pandas.pydata.org/docs/>
- [3] Scikit-learn documentation. Follow: <https://scikit-learn.org/stable/>
- [4] Matplotlib documentation. Follow: <https://matplotlib.org/stable/index.html>
- [5] Jupyter Notebook documentation. Follow: <https://docs.jupyter.org/en/latest/>
- [6] SMOTE for Imbalanced Classification with Python. Follow: <https://machinelearningmastery.com/smote-oversampling-for-imbalanced-classification/>
- [7] Decision tree – Scikit Learn. Follow: <https://scikit-learn.org/stable/modules/tree.html>
- [8] Random Forest Classifier. Follow: <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestClassifier.html>
- [9] XGBOOST documentation: Follow: <https://xgboost.readthedocs.io/en/stable/>
- [10] LightGBM documentation: Follow: <https://lightgbm.readthedocs.io/en/stable/>
- [11] catBoost documentation. Follow: <https://catboost.ai/>